

LLM - Prompt Injection Firewall

Sanjeev Ragunathan (sr2432)

INFO 5940, Prof. Jacob A. Matthews - Final Project

The Problem

- Prompt Injection, changes model behaviour
- Users paste sensitive information to the LLM (PII - SSN, Passport, Licenses, Keys, etc)
- The LLM providing sensitive/personal/secret data present in its context (info. present in training data and documents)


#1 OWASP LLM risk

Cisco - **30%** GenAI users enter personal/confidential information


GPT-4 has **31% ASR** on BIPiA indirect injection benchmark

The Threat Models


- Direct injection - the attacker is the user



- Indirect injection - the attacker plants payload in data LLM reads later



Our Scope



Out of scope, by design:


harmful-content requests (LLM05). The model's RLHF refuses those reliably - duplicating that work adds latency and false positives without benefit

Existing Defenses

System	Strength	Weakness
Llama Guard 3	String detection, 1B params	Slow, generative
Prompt Guard 2	Cheap, fast	Over-Defense
NeMo Guardrails	Programmable Rules	Rules don't generalize
StruQ	Removes the problem	Needs retraining

- Over-Defense introduced during Prompt Guard 2 in the InjecGuard paper 2024
- It's when the model identifies benign prompts as attacks - False Positives
- **Increased False Positives, reduces utility**
- Any prompts with "ignore" gets blocked.

Our Architecture




Evaluation

6-track evaluation; one dataset per threat

Track	Source	N
Direct Injection	deepset/prompt-injections	60
Indirect Injection	BIPIA (microsoft)	200
Inbound PII	ai4privacy/pii-masking-200k	200
Benign Control (FPR)	Dolly-15k	200
Sys-prompt Leak	Hand crafted - Lukas '23	15
PII echo leak	Hand crafted - Lukas '23	15

Results



- 100% leakage handling (system prompt + PII echo probes)
- 0% benign FPR (Dolly-15k control set)
- 26% BIPIA handling (hardest remaining threat)

Limitations & Future Work

What worked

- 100% handling on outbound leakage probes
- 0% false positives on Dolly-15k
- Strong PII handling with block/redact policies

Main limitation

- Indirect injection handled only 26%
- Attack is hidden inside long RAG/tool context
- Sentence-level classifiers lose the signal in noise

Future work

- StruQ-style separation of trusted vs. untrusted text
- Known-secrets list for system prompt / operator data
- Outbound classifier for paraphrased leakage
- Multilingual + international PII coverage

References

Articles:

- OWASP Foundation. OWASP Top 10 for LLM Applications, 2025.
- Greshake et al. Not What You've Signed Up For: Indirect Prompt Injection, 2023.
- Liu et al. Prompt Injection Attack Against LLM-Integrated Applications, 2024.
- Yi et al. BIPIA: Benchmarking Indirect Prompt Injection Attacks, 2023.
- Lukas et al. Analyzing Leakage of Personally Identifiable Information in Language Models, IEEE S&P 2023.
- Li et al. InjecGuard: Benchmarking and Mitigating Over-defense, 2024.
- Chen et al. StruQ: Defending Against Prompt Injection with Structured Queries, 2024.
- Chen et al. SecAlign: Aligning LLMs to Be Robust Against Prompt Injection, 2024.
- <https://www.arthur.ai/blog/from-jailbreaks-to-gibberish-understanding-the-different-types-of-prompt-injections>

Datasets:

deepset/prompt-injections, ai4privacy/pii-masking-200k, databricks/dolly-15k, BIPIA

Repository:

CheesyWasp Implementation - <https://www.github.com/sanjeev-ragunathan/llm-firewall>